

machine model

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [3]: df_1=pd.read_csv(r"C:\Users\arumu\Downloads\cardio_train.csv",delimiter=';')
df_1
```

```
Out[3]:
```

	id	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	alco	active	cardio
0	0	18393	2	168	62.0	110	80	1	1	0	0	1	0
1	1	20228	1	156	85.0	140	90	3	1	0	0	1	1
2	2	18857	1	165	64.0	130	70	3	1	0	0	0	1
3	3	17623	2	169	82.0	150	100	1	1	0	0	1	1
4	4	17474	1	156	56.0	100	60	1	1	0	0	0	0
...
69995	99993	19240	2	168	76.0	120	80	1	1	1	0	1	0
69996	99995	22601	1	158	126.0	140	90	2	2	0	0	1	1
69997	99996	19066	2	183	105.0	180	90	3	1	0	1	0	1
69998	99998	22431	1	163	72.0	135	80	1	2	0	0	0	1
69999	99999	20540	1	170	72.0	120	80	2	1	0	0	1	0

70000 rows × 13 columns

```
In [4]: # Display the first few rows of the dataset
df_1.head()

# Check for missing values
print(df_1.isnull().sum())

# Basic information about the dataset
print(df_1.info())
```

```

id          0
age         0
gender      0
height      0
weight      0
ap_hi       0
ap_lo       0
cholesterol 0
gluc        0
smoke       0
alco        0
active      0
cardio      0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 70000 entries, 0 to 69999
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              70000 non-null  int64
1   age             70000 non-null  int64
2   gender          70000 non-null  int64
3   height          70000 non-null  int64
4   weight          70000 non-null  float64
5   ap_hi           70000 non-null  int64
6   ap_lo           70000 non-null  int64
7   cholesterol     70000 non-null  int64
8   gluc            70000 non-null  int64
9   smoke           70000 non-null  int64
10  alco            70000 non-null  int64
11  active          70000 non-null  int64
12  cardio          70000 non-null  int64
dtypes: float64(1), int64(12)
memory usage: 6.9 MB
None

```

```

In [5]: # Split dataset into features (X) and target (y)
X = df_1[['age', 'height', 'weight', 'ap_hi', 'ap_lo', 'cholesterol', 'gluc', 'smoke', 'alco', 'active']]
y = df_1['cardio']
X
y

```

```

Out[5]: 0      0
1      1
2      1
3      1
4      0
..
69995  0
69996  1
69997  1
69998  1
69999  0
Name: cardio, Length: 70000, dtype: int64

```

```

In [6]: # Ensure there are no categorical columns that need to be encoded
print(X.dtypes)

```

```

age          int64
height       int64
weight       float64
ap_hi        int64
ap_lo        int64
cholesterol  int64
gluc         int64
smoke        int64
alco         int64
active       int64
dtype: object

```

```

In [7]: # Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
X_train, X_test, y_train, y_test

```

```
Out[7]: (
      age  height  weight  ap_hi  ap_lo  cholesterol  gluc  smoke  alco  \
68681  20417    160    64.0   120    90             3     1     0     0
19961  22690    167    65.0   120    80             3     3     0     0
11040  22784    160    66.0   120    90             1     1     0     0
27673  22648    163    55.0   125    90             3     1     0     0
22876  21712    158    85.0   150    80             3     1     0     0
...      ...      ...      ...      ...      ...      ...      ...      ...      ...
37194  16001    170    75.0   150    80             1     1     1     0
6265   23209    162    73.0   160    90             1     1     0     0
54886  23589    169    74.0   120    80             1     1     0     0
860    18227    167    70.0   120    80             1     1     0     0
15795  15114    177    64.0   120    80             1     1     0     0
```

```

      active
68681      1
19961      0
11040      1
27673      1
22876      1
...      ...
37194      1
6265      1
54886      1
860       0
15795      1
```

```
[49000 rows x 10 columns],
      age  height  weight  ap_hi  ap_lo  cholesterol  gluc  smoke  alco  \
46730  21770    156    64.0   140    80             2     1     0     0
48393  21876    170    85.0   160    90             1     1     0     0
41416  23270    151    90.0   130    80             1     1     0     0
34506  19741    159    97.0   120    80             1     1     0     0
43725  18395    164    68.0   120    80             1     1     0     0
...      ...      ...      ...      ...      ...      ...      ...      ...      ...
1216   22392    161    68.0   150   100             2     1     0     0
19036  14462    168    66.0   130    80             1     1     0     0
51256  14805    159    81.0   130   100             1     1     0     0
48198  20519    143    65.0   130    90             1     1     0     0
2571   16181    156    80.0   180   100             2     1     0     0
```

```

      active
46730      1
48393      1
41416      1
34506      1
43725      1
...      ...
1216      1
19036      1
51256      0
48198      1
2571      1
```

```
[21000 rows x 10 columns],
68681      1
19961      0
11040      1
27673      1
22876      1
...
37194      1
6265      1
54886      0
860       0
15795      0
Name: cardio, Length: 49000, dtype: int64,
46730      1
48393      1
41416      1
34506      1
43725      0
...
1216      1
19036      0
51256      0
48198      1
2571      1
Name: cardio, Length: 21000, dtype: int64)
```

```
In [8]: # Standardize the data (important for Random Forest as well)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

X_test

```
Out[8]: array([[ 0.93597822, -1.01890093, -0.70816849, ..., -0.31319072,
               -0.24186407,  0.49466891],
               [ 0.97889556,  0.68916043,  0.75248336, ..., -0.31319072,
               -0.24186407,  0.49466891],
               [ 1.54329899, -1.62892285,  1.10025762, ..., -0.31319072,
               -0.24186407,  0.49466891],
               ...,
               [-1.88401453, -0.65288778,  0.47426396, ..., -0.31319072,
               -0.24186407, -2.02155415],
               [ 0.4294727 , -2.60495792, -0.63861364, ..., -0.31319072,
               -0.24186407,  0.49466891],
               [-1.32689895, -1.01890093,  0.40470911, ..., -0.31319072,
               -0.24186407,  0.49466891]])
```

```
In [9]: # Initialize and train the Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
```

```
Out[9]: ▼ RandomForestClassifier ⓘ ?
RandomForestClassifier(random_state=42)
```

```
In [10]: # Predict the test set
y_pred = rf_model.predict(X_test)

# Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Random Forest Accuracy: {accuracy * 100:.2f}%")

# Print confusion matrix and classification report
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Random Forest Accuracy: 71.34%

Confusion Matrix:

```
[[7568 2893]
 [3126 7413]]
```

```
In [11]: print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```


Classification Report:

	precision	recall	f1-score	support
0	0.71	0.72	0.72	10461
1	0.72	0.70	0.71	10539
accuracy			0.71	21000
macro avg	0.71	0.71	0.71	21000
weighted avg	0.71	0.71	0.71	21000

```
In [12]: feature_importances = pd.Series(rf_model.feature_importances_, index=X.columns)
```

```
In [15]: # Feature importance
plt.figure(figsize=(10,6))
feature_importances.sort_values(ascending=False).plot(kind='bar')
plt.title('Feature Importances for Heart Disease Detection using Random Forest')
plt.ylabel('Importance')
plt.show()
```

Feature Importances for Heart Disease Detection using Random Forest



In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js